

A Depth-Aware, Physics-Based Data Augmentation Framework for Simulating Rain and Fog Distortions in Outdoor-**RGB Images**

Rongxuan Duan
Pioneer Research Institute

Abstract

1 Introduction

Images used in computer vision are often affected by distortions. Distortion in Computer Vision is defined as “any signal that produces a deviation in the visual perception of the image/video from its original expectation is called distortion. But when the original image or video is unavailable, any component of the observed signal presenting a level of visual discomfort or simply annoying effect on the aesthetic visual appearance is considered as a distortion” [1]. These distortions could be sorted as In-Capture Distortion, including blur, noise, uneven illumination, color shifts, haze, rain, snow, and other visibility changes; and Post-Capture Distortion, including transmission, impairments, enhancement artefacts, and adversarial distortion. Although some distortions may appear minor to human observers, they can change the statistical structure of an image and reduce the reliability of machine learning models trained mainly on clean data [1]. Distortions affect the performances of high-level tasks, such as object classification, object detection, and scene understanding [1]. Therefore, distortion detection and removal are a possible preprocessing step for high-level tasks to decrease the impact of distortion on performance.

For machine learning, distortion is not only an image-quality issue but also a dataset and training issue. Models usually learn from the distribution of inputs they are given, so a model trained only on clean images may have poor generalizability when working on the same object appearing under poor visibility, weather effects, or degraded contrast. Distortion-specified datasets are therefore important because they make it possible to evaluate model robustness by distortion type and severity, rather than reporting only a single clean-data accuracy score [1].

Current datasets of natural distortions still have important limitations. Real rain, haze, fog, and snow images can be difficult to collect at scale, and the resulting data may be unbalanced across weather type, severity, scene cat-

egory, and annotation quality. Many benchmark datasets also include distortions as broad corruptions rather than as physically meaningful, scene-dependent effects. This limits how well they can support controlled robustness studies or targeted training strategies. The need for environment-aware image formation is also visible in underwater computer vision, where more complicated distortions also present due to an uneven refractive index of water, such as attenuation, scattering, color-shifting, halo effect, etc. [2].

Related work has explored both distortion recognition and natural-distortion simulation. Survey work on distortion classification emphasizes that multiple and natural distortions remain difficult to categorize and model consistently [1]. In weather-focused vision, rain simulation has been used to render controlled rain effects for robustness evaluation, while haze and fog modeling often relies on atmospheric image-formation ideas such as scattering, attenuation, and reduced scene contrast. Work on adverse-weather perception, including foggy driving scenes, also shows that real weather can strongly affect downstream recognition and detection performance.

Our work trains a distortion augmentation framework, which adds controlled rain and fog effects to existing baseline datasets such as . The framework is evaluated by comparing the performances of downstream object detection models, such as , on the original dataset and the augmented dataset.

2 Related Work

2.1 Distortion Classification and Evaluation

Traditional distortion classification models and evaluation use blur indices, noise estimators, natural scene statistics, image quality metrics, Support Vector Machines (SVM), and K-Nearest Neighbours (KNN) for evaluating the impact of distortions. Recent distortion classification and detection

Commented [PW1]: Hi Rongxuan!
I enjoyed reading your paper, it's very informative and you did a nice job in explaining the types of distortion and their impact. My notes focus on word-choice and some grammatical agreements in plural/singular.
Good luck on the next draft!
-Auden

Commented [PW2]: "...are often distorted." Always try to make your writing as concise and clear as possible.

Commented [PW3]: If this should be capitalized, then it also needs to be in the previous sentence. If not, de-capitalize this one. Stay consistent throughout the paper

Commented [PW14]: cut

Commented [PW4]: I recommend integrating this quote better. As is, it's a very long, multi-sentence quote which loses the "punchiness" a quote is supposed to provide. Quotes aren't for whole definitions, they are supposed to be clips of the most important, supportive evidence of your statement/claim

Commented [PW5]: since you just used "including" try to find another synonym here

Commented [PW6]: There's a lot of "distortion" mention. If you can't find a synonym, I think simply saying "these effects" or "these changes" would work here

Commented [PW7]: primarily*

Commented [PW15]: In literature it's customary to avoid "I" "We" "Our" terms, instead using passive language, making the sentence "The model trains a distortion augmentation framework" etc

Commented [PW8]: You can cut this since you previously mentioned distortion in the same sentence, so the reader will know to assume the subject hasn't changed

Commented [PW16]: such as?

Commented [PW9]: Try to avoid repetition, as some of my previous comments have stated. Here you repeat issue twice. Instead, you can rewrite the sentence to say something akin to "image-quality and dataset training issues" or just another synonym: "problems" "concerns" "setbacks" etc.

Commented [PW10]: cut

Commented [PW17]: ?

Commented [PW11]: cut

Commented [PW12]: similar to my previous comment: you mention "distortion-specified datasets" already, which the reader will refer back to. So you can cut this second "distortion"

Commented [PW13]: Try reading this sentence with a without the "still." With it, it sounds conversational. Without it, the sentence grabs the reader's attention much better -- this is crucial in literature and writing, making it as easy as possible for the reader to absorb the information you're discussing

Commented [PW18]: to evaluate

models use Convolutional Neural Networks (CNN) and temporal models. Ahn, Kang, and Sohn train a Y-shaped CNN for joint distortion classification and localization. Other methods for distortion classification include feature pool, relating metrics with specific distortions, and Natural Scene Statistics (NSS). The method supports learned distortion features, yet most datasets focus on controlled distortions rather than natural weather [1, 3].

The quality of visual content in images and distortion severity are important factors that affect performances in high-level tasks, such as object classification, object detection, and scene understanding. Using distortion classification and detection models could help researchers finetune models that are more resistant to distortions. However, our work helps finetuning models by presenting a data augmentation framework that increases models' resilience to rain and fog distortions from training, which does not require distortion detection and distortion removal process.

2.2 Current Limitations of Distortion Datasets

Current distortion datasets give labels and severity levels, but many examples come from synthetic, single-type, quality-assessment settings. LIVE, TID2008, TID2013, CSIQ, Waterloo, KADID-10K, LIVEMD, and MDID help controlled evaluation. They do not give natural weather with precise physical labels. Rain, haze, fog, and snow depend on lighting, exposure, particle density, scene depth, and occlusion. There is a lack of labeled, controlled multi-distortion datasets. Consequently, classification and localization of multiple distortions remains challenging due to lack of labeled, controlled datasets [1, 4]. Our work adds controlled, artificial distortions, which addresses lack of labeled, controlled datasets, especially for outdoor datasets.

2.3 Traditional Rain Rendering

Garg and Nayar model rain streak photometry and spatio-temporal behavior. Their work links rain visibility to drop size, velocity, exposure time, scene radiance, camera settings, and drop shape oscillation. This result matters for rendering because a transparent 2D texture ignores camera and scene factors. Traditional rain rendering lacks the effect of moisture on the scene, and it often appears to be unrealistic. For instance, traditional rain rendering lacks wet surfaces such as water on the ground, and it creates raindrops that are supposed to be hidden behind buildings [4]. Though, this framework supported rain detection, removal, rendering, and camera-based rain measurement [5].

2.4 Modern Rain Rendering

Modern rain rendering turns rain models into dataset augmentation for robustness tests. Halder, Lalonde, and de Charette combine a particle simulator, outdoor illumination estimation, and rain photometry to render rain and fog over clear driving images. The pipeline preserves labels from KITTI and Cityscapes while changing weather appearance. Detection and segmentation results then measure model sensitivity under controlled rain [6].

Tremblay, Halder, de Charette, and Lalonde compare three rendering routes: physics-based, data-driven, and hybrid. They classify raindrops into rain particles and fog-like rain. Fog-like rain is a single pixel that contains multiple rain particles; it locates far away from the camera and appears as fog. The physics-based route combines fog-like rain, rain particles, and photometric modeling. The renderer estimates outdoor illumination because a raindrop receives light from a wider angular range (approximately 160 degree) than the camera view (70 to 100 degree). This design improves physical control over a flat overlay. The same design also depends on sky visibility, environmental symmetry about photometry, and compute cost [4, 7, 8].

The data-driven route uses a GAN-style image translation model. A GAN trains two networks: a generator synthesizes target-domain images, and a discriminator judges whether each image looks real or generated. Adversarial training on the generator outputs rainy-looking images. The output from the generator contains feature such as dark sky, rainy environment, wet surfaces and dim lamination. However, GAN is less controlled, which risks for rain density, depth ordering, and scene consistency. GAN's focus on global appearance also creates unrealistic features, such as fake water reflections [4, 9, 10].

Tremblay et al. render rain over KITTI, Cityscapes, and nuScenes. They evaluate object detection, semantic segmentation, and depth estimation under increasing degradation. The paper reports degrading performances as increasing rain intensity and more robust performances after finetuning using rain-augmented datasets. The paper reports that outputs from GAN only are user-rated as the most realistic images; outputs from GAN plus PBR perform the best on data augmentation [4].

Our work is inspired from the framework from Tremblay et al. our work uses the traditional rain rendering framework with depth estimation to build rain streaks with controlled, depth-dependent rain streaks; our work also involves using a Cycle-GAN model that adds rainy and foggy visual effects.

2.5 Fog Rendering

Fog rendering often follows the atmospheric image formation model. Scene radiance attenuates with distance, while airlight grows along the camera ray. Sakaridis, Dai, and

Commented [PW19]: performance, singular

Commented [PW20]: including

Commented [PW23]: is located*

Commented [PW24]: degrees*

Commented [PW25]: camera's* view (possessive)

Commented [PW26]: Always plural "degrees" when referencing weather temp

Commented [PW21]: this sentence isn't grammatically correct. Either "help keep a controlled evaluation" or "help control the evaluation"

Commented [PW22]: This is worded in a way that's difficult to understand. Can you rephrase it? Maybe produce* rather than "give"

Commented [PW27]: features* (make the plural agreement)

Commented [PW28]: cut "for"

Commented [PW29]: camera's* (possessive)

Van Gool apply this model to clear Cityscapes images and produce Foggy Cityscapes. Depth makes fog stronger for far regions rather than uniform across the image. Training with synthetic fog improves semantic foggy scene understanding, but results depend on depth quality, density choices, and the synthetic-to-real gap [11].

Clear2Fog extends this direction with monocular depth estimation and atmospheric light estimation. Monocular depth estimation does not require 3D constructions such as LiDAR; rather, this paper uses Apple Depth Pro to produce an estimated depth map. The pipeline simulates fog on clear-weather driving datasets and tries to preserve camera-LiDAR consistency. Without relying on LiDAR, Clear2Fog remains stable under extreme conditions like fog, snow or haze. Seeing Through Fog Without Seeing Fog focuses on multimodal perception rather than rendering. The dataset and fusion model show different failure patterns across camera, LiDAR, radar, and gated near-infrared sensors [12, 13].

2.6 Depth Estimation and Application in Distortion Rendering

Depth provides the geometry needed for natural distortion placement. In haze and fog, depth sets attenuation and airlight strength. In rain, depth helps determining the severity, intensity and placement of rain streaks. A renderer without geometry often places plausible-looking weather in physically inconsistent locations [11, 4].

Depth estimation models estimate distance from an image or complete sparse sensor depth. The output supplies an approximate geometry map when a dataset lacks dense ground truth. Jaritz et al. study sparse and dense depth completion for driving scenes. Depth Anything V2 offers a modern monocular option trained with synthetic images and pseudo-labeled real images. This makes Depth Anything V2 useful for RGB-only distortion rendering [14, 15].

Our work uses the depth-map output from Depth Anything V2 as the guiding parameter for transparency and appearance of rain streaks and translucency and intensity of fog. The depth-map output from Depth Anything V2 is stable and reliable, and it is a strong model for completing our computer vision tasks.

2.7 Datasets With Distortions

General distortion and corruption datasets support baselines, but natural weather rendering requires more geometry and scene context. The distortion survey lists LIVE, TID2008, TID2013, CSIQ, Waterloo, KADID-10K, LIVEMD, and MDID as major image-quality or multi-distortion resources. ImageNet-C standardizes common corruption robustness for image classifiers. These resources help robustness framing. They do not model depth-aware rain, haze, or fog. [1].

Weather-specific datasets match the project more closely.

ACDC includes fog, nighttime, rain, and snow with dense driving-scene annotations and normal-condition correspondences for many scenes. Foggy Cityscapes and rain-augmented KITTI, Cityscapes, and nuScenes preserve existing labels while adding synthetic weather. The limitation of these datasets is the lack of controlled severity or intensity, depth-aware labels, and labeled natural distortions [16, 11, 4].

2.8 BDD100K

BDD100K serves as a baseline dataset for our work, because the data are large, diverse, and driving-focused. The dataset contains 100,000 driving videos and supports detection, segmentation, lane marking, drivable area, and related driving labels. BDD100K also includes geographic, environmental, and weather variation. The limitation is missing paired clean/adverse views and controlled severity labels. A depth-aware rain and haze renderer would turn BDD100K into a controlled robustness testbed while keeping real driving diversity. [17].

References

- [1] Muhammad Ali Qureshi et al. *Distortion Classification in Computer Vision Applications: Current Progress, Challenges, and Perspectives — ACM Computing Surveys*. <https://dl.acm.org/doi/10.1145/3773023>. (Visited on 06/10/2026).
- [2] *A Survey on Underwater Computer Vision — ACM Computing Surveys*. <https://dl.acm.org/doi/10.1145/3578516>. (Visited on 06/10/2026).
- [3] Namhyuk Ahn, Byungkong Kang, and Kyung-Ah Sohn. “Image Distortion Detection Using Convolutional Neural Network”. In: *2017 4th IAPR Asian Conference on Pattern Recognition (ACPR)*. Nanjing, China: IEEE, Nov. 2017, pp. 220–225. ISBN: 978-1-5386-3354-0. DOI: 10.1109 / ACPR . 2017 . 95. (Visited on 06/10/2026).
- [4] Maxime Tremblay et al. “Rain Rendering for Evaluating and Improving Robustness to Bad Weather”. In: *International Journal of Computer Vision* 129.2 (Feb. 2021), pp. 341–360. ISSN: 1573-1405. DOI: 10.1007 / s11263 - 020 - 01366 - 3. (Visited on 06/10/2026).
- [5] Kshitiz Garg and Shree K. Nayar. “Vision and Rain”. In: *International Journal of Computer Vision* 75.1 (Oct. 2007), pp. 3–27. ISSN: 1573-1405. DOI: 10.1007 / s11263 - 006 - 0028 - 6. (Visited on 06/14/2026).

Commented [PW30]: applies*

Commented [PW32]: try to find another synonym here to avoid repetition

Commented [PW31]: determine*

- [6] Shirsendu Sukanta Halder, Jean-Francois Lalonde, and Raoul de Charette. *Physics-Based Rendering for Improving Robustness to Rain*. Aug. 2019. DOI: 10.48550/arXiv.1908.10335. arXiv: 1908.10335 [cs.CV]. (Visited on 06/14/2026).
- [7] Yannick Hold-Geoffroy et al. *Deep Outdoor Illumination Estimation*. Apr. 2018. DOI: 10.48550/arXiv.1611.06403. arXiv: 1611.06403 [cs.CV]. (Visited on 06/14/2026).
- [8] Jinsong Zhang et al. *All-Weather Deep Outdoor Lighting Estimation*. June 2019. DOI: 10.48550/arXiv.1906.04909. arXiv: 1906.04909 [cs.CV]. (Visited on 06/16/2026).
- [9] Phillip Isola et al. *Image-to-Image Translation with Conditional Adversarial Networks*. Nov. 2018. DOI: 10.48550/arXiv.1611.07004. arXiv: 1611.07004 [cs.CV]. (Visited on 06/16/2026).
- [10] Xun Huang et al. "Multimodal Unsupervised Image-to-Image Translation". In: *Computer Vision – ECCV 2018*. Ed. by Vittorio Ferrari et al. Cham: Springer International Publishing, 2018, pp. 179–196. ISBN: 978-3-030-01219-9. DOI: 10.1007/978-3-030-01219-9_11.
- [11] Christos Sakaridis, Dengxin Dai, and Luc Van Gool. "Semantic Foggy Scene Understanding with Synthetic Data". In: *International Journal of Computer Vision* 126.9 (Sept. 2018), pp. 973–992. ISSN: 1573-1405. DOI: 10.1007/s11263-018-1072-8. (Visited on 06/14/2026).
- [12] Mohamed Ahmed Mohamed and Xiaowei Huang. *A Data Efficiency Study of Synthetic Fog for Object Detection Using the Clear2Fog Pipeline*. May 2026. DOI: 10.48550/arXiv.2605.12608. arXiv: 2605.12608 [cs.CV]. (Visited on 06/18/2026).
- [13] Mario Bijelic et al. "Seeing Through Fog Without Seeing Fog: Deep Multimodal Sensor Fusion in Unseen Adverse Weather". In: *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*. June 2020, pp. 11682–11692. DOI: 10.1109/CVPR42600.2020.01170.
- [14] Maximilian Jaritz et al. *Sparse and Dense Data with CNNs: Depth Completion and Semantic Segmentation*. Aug. 2018. DOI: 10.48550/arXiv.1808.00769. arXiv: 1808.00769 [cs.CV]. (Visited on 06/14/2026).
- [15] Lihe Yang et al. *Depth Anything V2*. <https://arxiv.org/abs/2406.09414v2>. June 2024. (Visited on 06/24/2026).
- [16] Christos Sakaridis et al. *ACDC: The Adverse Conditions Dataset with Correspondences for Robust Semantic Driving Scene Perception*. <https://arxiv.org/abs/2104.13395v5>. Apr. 2021. (Visited on 06/24/2026).
- [17] Fisher Yu et al. *BDD100K: A Diverse Driving Dataset for Heterogeneous Multitask Learning*. <https://arxiv.org/abs/1805.04687v2>. May 2018. (Visited on 06/24/2026).
- [18] Md. Imtiyaz Anwar and Arun Khosla. "Vision Enhancement through Single Image Fog Removal". In: *Engineering Science and Technology, an International Journal* 20.3 (June 2017), pp. 1075–1083. ISSN: 22150986. DOI: 10.1016/j.jestech.2016.11.015. (Visited on 06/18/2026).
- [19] Mohamed Ahmed Mohamed and Xiaowei Huang. *A Data Efficiency Study of Synthetic Fog for Object Detection Using the Clear2Fog Pipeline*. <https://arxiv.org/abs/2605.12608v1>. May 2026. (Visited on 06/24/2026).
- [20] C. K. Wolfe. *BDD100K: A Large-scale Diverse Driving Video Database*. (Visited on 06/14/2026).